Report of Pneumonia images's classification

Made by:

- CHAGOUTI El houcine
- DARBALI Oussama

Under the guidance of:

- MAHMOUDI Abdelhak

Table of contents:

I- What is Pneumonia?

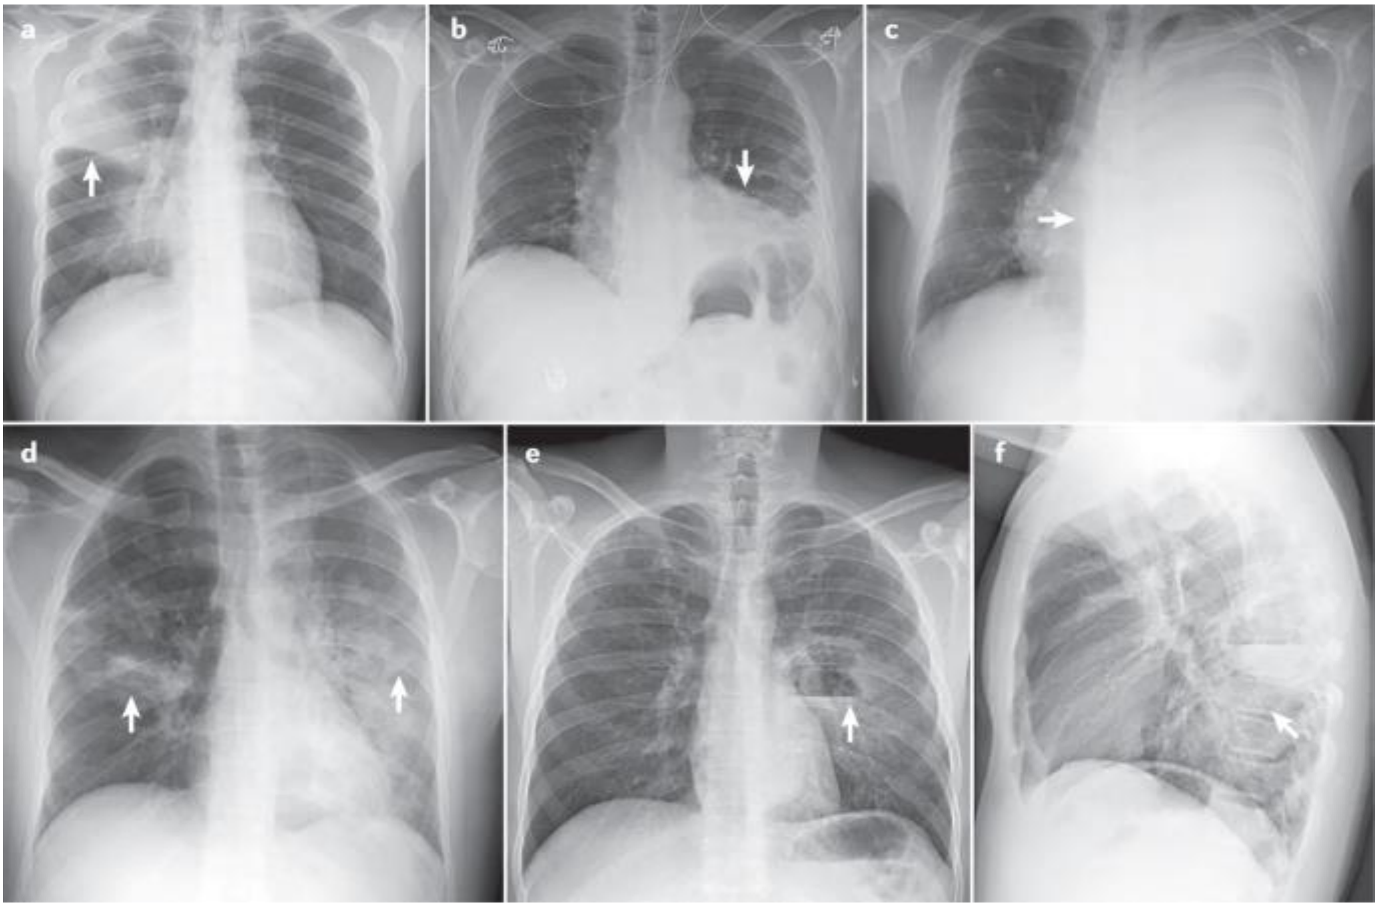

Pneumonia, a prevalent respiratory infection, encompasses a range of inflammatory conditions affecting the lungs, primarily caused by bacterial, viral, or fungal agents. Its clinical manifestations vary from mild to severe, presenting symptoms such as cough, fever, and difficulty breathing. Pneumonia poses a significant public health burden worldwide, particularly impacting vulnerable populations like children, the elderly, and individuals with compromised immune systems. Timely and accurate diagnosis is crucial for effective management and treatment, as delayed or inadequate intervention can lead to complications and even mortality. Chest X-ray imaging plays a pivotal role in pneumonia diagnosis, offering insights into lung abnormalities indicative of the disease. Leveraging advanced technologies like deep learning for pneumonia detection holds promise in enhancing diagnostic efficiency, aiding clinicians in prompt decision-making, and ultimately improving patient outcomes.

II- Why building Models to predict Pneumonia?

The development of deep learning models for pneumonia prediction addresses a pressing need in healthcare by revolutionizing diagnostic capabilities. Traditional pneumonia diagnosis often relies on human interpretation of medical imaging, which can be time-consuming, subjective, and prone to error. By harnessing the power of deep learning algorithms, healthcare providers can streamline the diagnostic process, expediting the identification of pneumonia cases and facilitating early intervention. Moreover, the scalability and reproducibility of deep learning models offer the potential for widespread implementation across healthcare settings, including resource-limited regions where access to skilled radiologists may be limited. By automating pneumonia detection from chest X-ray images, deep learning models not only enhance diagnostic accuracy but also optimize healthcare resource allocation and improve patient outcomes. As such, investing in the development of deep learning models for pneumonia prediction represents a significant advancement in medical technology with far-reaching implications for global healthcare delivery.

III- Used Technologies:

1- Tensorflow and Keras:

TensorFlow and Keras are two fundamental tools in the realm of deep learning, playing pivotal roles in model development, training, and deployment. TensorFlow, an open-source machine learning framework developed by Google, provides a comprehensive ecosystem for building and deploying machine learning models, including deep neural networks. Its flexible architecture supports a range of deployment options, from desktop to cloud and mobile devices, making it versatile for various applications. Keras, on the other hand, is a high-level neural networks API, integrated seamlessly with TensorFlow, simplifying the process of building and experimenting with deep learning models. With its user-friendly interface and intuitive design, Keras enables rapid prototyping of neural networks, allowing researchers and developers to focus on model architecture and experimentation rather than low-level implementation details. Together, TensorFlow and Keras empower practitioners to leverage the full potential of deep learning, facilitating innovation and advancements across diverse domains, from healthcare to finance and beyond.

2- Streamlit:

Streamlit

Streamlit, a Python library for building interactive web applications, has emerged as a valuable tool for deploying and showcasing deep learning models to a broader audience. With Streamlit, developers can create user-friendly interfaces for exploring and interacting with machine learning applications, including those powered by TensorFlow and Keras. By leveraging Streamlit's intuitive APIs and customizable components, developers can quickly turn their deep learning models into interactive web apps without the need for extensive web development experience. Streamlit's real-time updates and sharing capabilities further enhance collaboration and dissemination of deep learning projects, enabling stakeholders to engage with models and insights in a seamless and accessible manner. Whether it's visualizing model predictions, exploring data insights, or conducting parameter tuning, Streamlit empowers developers to create compelling and interactive experiences that democratize access to deep learning technology and foster innovation in the broader community.

IV- Used Techniques:

1- Simple CNN:

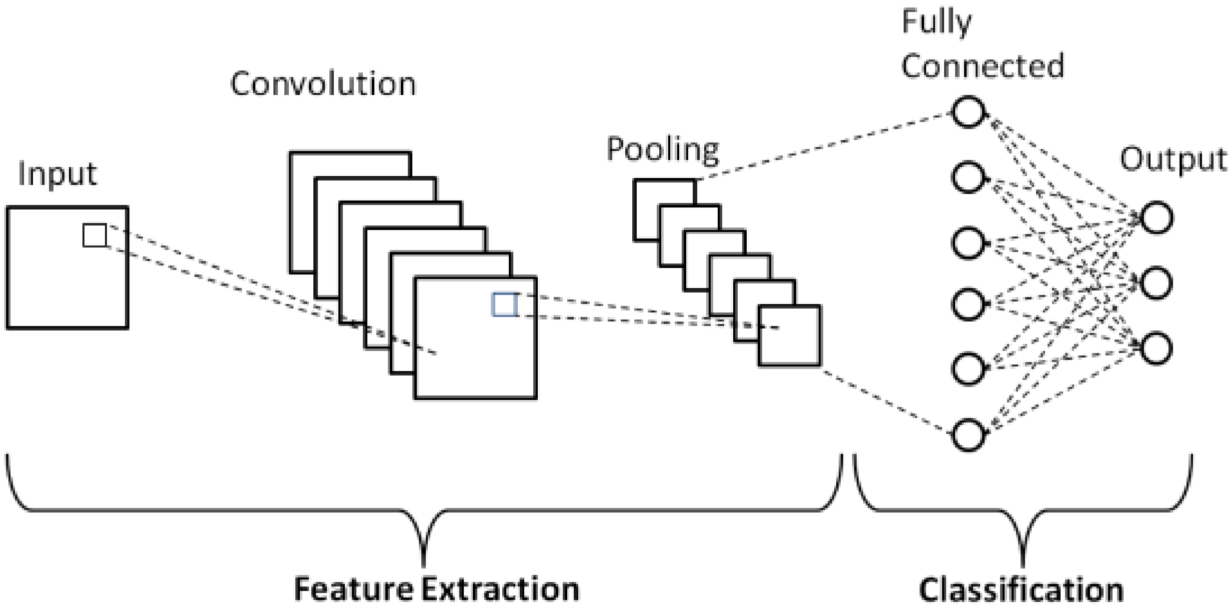

At first, we employed a simple Convolutional Neural Network (CNN) architecture for our pneumonia prediction task. The CNN comprises several key layers essential for image classification tasks. Initially, a convolutional layer with 32 filters of size 3x3 was added to the model, utilizing the ReLU activation function to introduce non-linearity and extract features from the input images, which were standardized to dimensions of 64x64 pixels with RGB channels. Subsequently, a max-pooling layer with a pool size of 2x2 was incorporated to downsample the feature maps, reducing computational complexity while preserving important features. Following this, a second convolutional layer identical to the first was introduced, enabling the extraction of deeper and more abstract features from the data. Another max-pooling layer followed to further condense the feature maps. The subsequent flattening layer reshaped the 2D feature maps into a 1D vector, preparing them for input into the densely connected layers. Two fully connected layers were then added, consisting of 128 neurons each, with ReLU activation in the first and sigmoid activation in the final layer, which served as the output layer for binary classification. Finally, the model was compiled using the Adam optimizer and binary cross-entropy loss function, with accuracy as the evaluation metric. This sequential arrangement of layers forms the backbone of our CNN model, facilitating the extraction and classification of relevant features from chest X-ray images to predict the presence of pneumonia.

The testing accuracy achieved by our CNN model is approximately 75.96%. While this accuracy demonstrates a reasonable performance, there may still be room for improvement. Several factors could contribute to the model's performance, including the complexity of the dataset, the adequacy of the model architecture, and the quality of data preprocessing. Fine-tuning hyperparameters, increasing the depth of the network, or implementing more advanced techniques such as transfer learning could potentially enhance the model's accuracy further. Additionally, exploring methods to mitigate issues like overfitting or class imbalance within the dataset could also lead to improved performance. Nonetheless, achieving a testing accuracy of 75.96% underscores the efficacy of the CNN model in predicting pneumonia from chest X-ray images and provides a solid foundation for future iterations and enhancements.

2- Data augmentation and CNN:

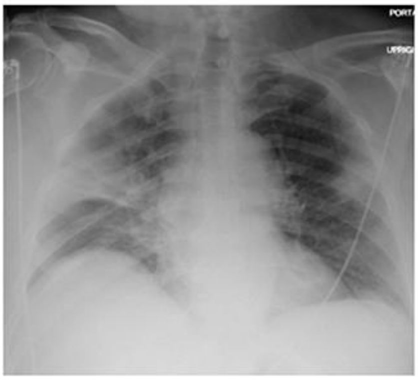

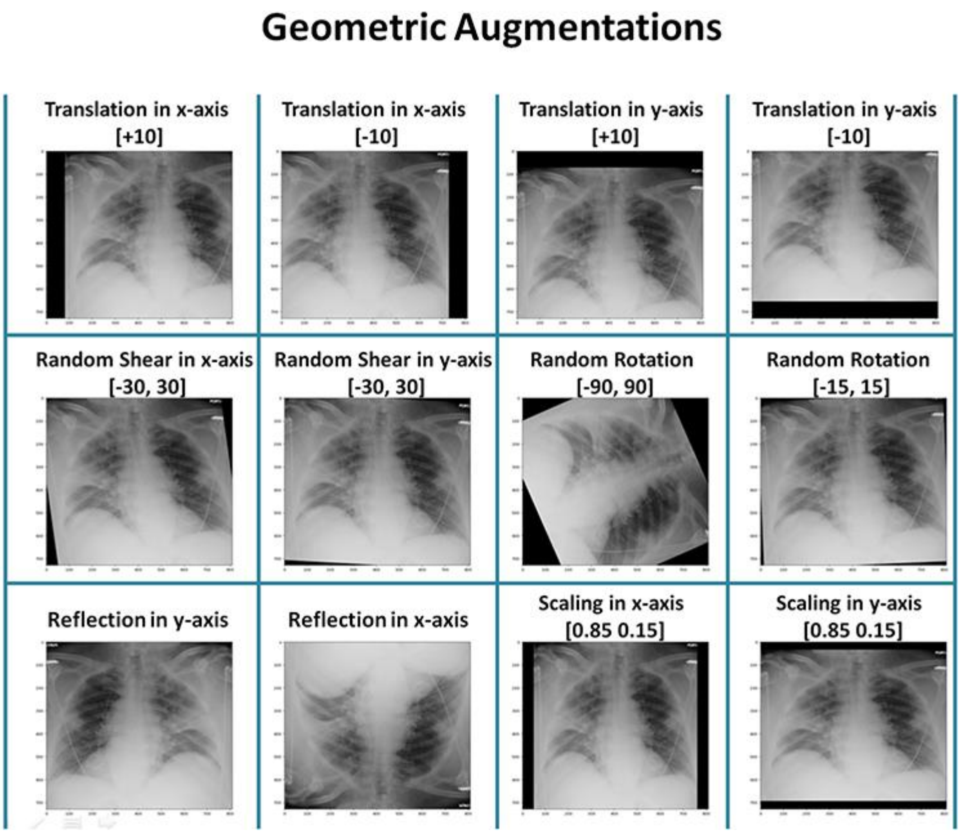

In this iteration of our pneumonia prediction project, we have employed data augmentation techniques along with a more complex Convolutional Neural Network (CNN) architecture. The data augmentation, facilitated by the ImageDataGenerator from Keras, introduces variations in the training dataset, thereby enhancing model generalization and robustness. By applying random rotations, zooms, shifts, and flips to the input images, we augment the training set to capture a wider range of variations and improve the model's ability to generalize to unseen data. The CNN architecture comprises multiple convolutional layers interspersed with batch normalization, dropout layers, and max-pooling layers. This deeper and more intricate network architecture enables the extraction of hierarchical features from the input images, capturing both low-level and high-level representations relevant to pneumonia detection. Specifically, the model consists of convolutional layers with increasing filter depths, followed by batch normalization to stabilize and accelerate the training process. Dropout layers are strategically inserted to mitigate overfitting by randomly deactivating neurons during training. The max-pooling layers downsample the feature maps, reducing computational complexity while preserving essential information. The fully connected layers at the end of the network integrate the extracted features for final classification. The model is compiled using the RMSprop optimizer with binary cross-entropy loss, and accuracy is chosen as the evaluation metric. This complex CNN architecture, combined with data augmentation, aims to further enhance the accuracy and robustness of pneumonia prediction from chest X-ray images, paving the way for more reliable diagnostic tools in healthcare settings.

The testing accuracy achieved by our enhanced CNN model, incorporating data augmentation techniques and a more complex architecture, is an impressive 91.67%. This substantial improvement from our previous model underscores the effectiveness of the strategies employed. The utilization of data augmentation has allowed the model to learn from a more diverse set of training examples, thus enhancing its ability to generalize to unseen data. Additionally, the incorporation of a deeper and more intricate CNN architecture enables the extraction of more nuanced features from the chest X-ray images, leading to better discrimination between pneumonia and non-pneumonia cases. The inclusion of batch normalization and dropout layers further contributes to the model's robustness by stabilizing and regularizing the training process, thereby mitigating overfitting. Overall, achieving a testing accuracy of 91.67% demonstrates the efficacy of our improved CNN model in accurately predicting pneumonia from medical images, highlighting its potential as a valuable tool in clinical diagnosis and decision-making.

3- Using Ensemble Learning Techniques:

3.1- CNN Bagging:

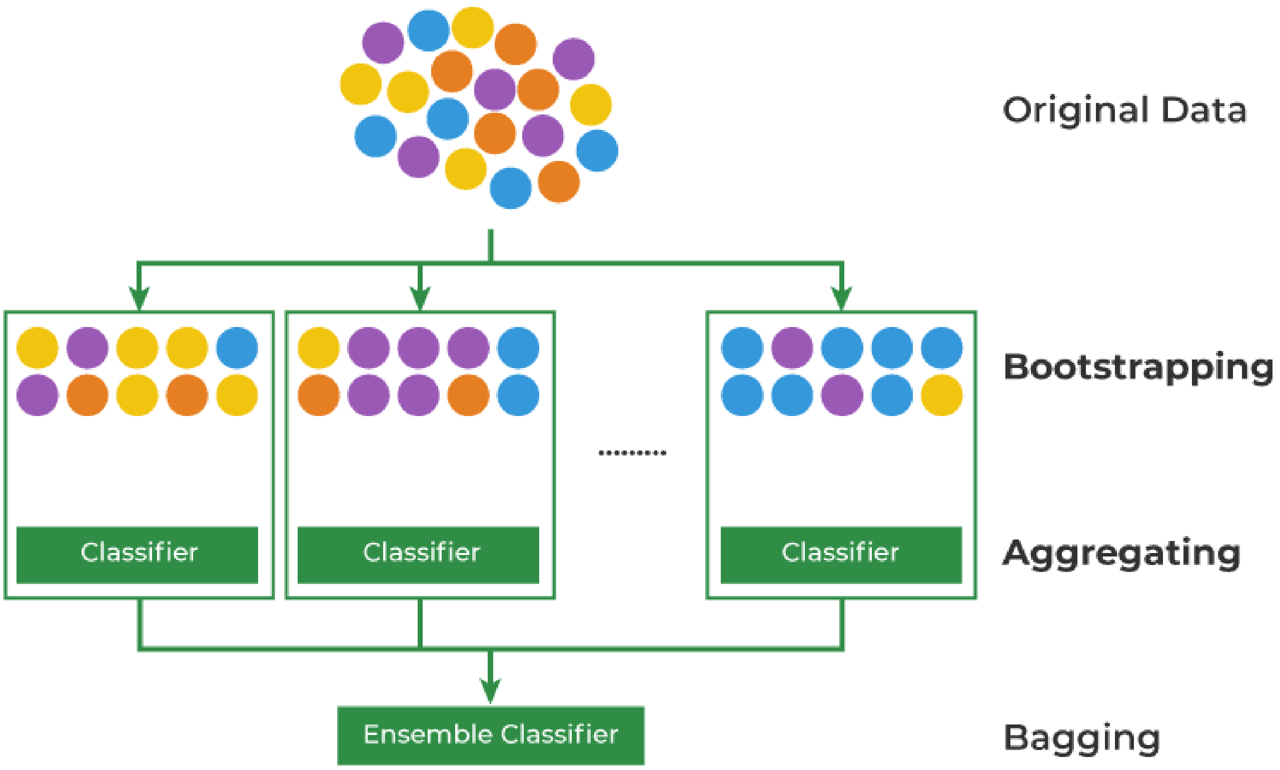

In our quest to enhance the performance of our simple CNN model for pneumonia prediction, we adopted the bagging technique, which involves training multiple models on different subsets of the training data and combining their predictions through a voting mechanism. To implement bagging, we first defined a base CNN model architecture within a function called build_cnn_model(), comprising convolutional, max-pooling, and dense layers for classification. We then instantiated five models using this function, each trained on a distinct subset of the training data. After training, we aggregated their predictions using a voting mechanism, where the final prediction is determined by a majority vote among the models. By leveraging the collective knowledge of these diverse models, bagging reduces overfitting and enhances the robustness of our pneumonia prediction system, aiming for higher accuracy and reliability in diagnosing pneumonia from chest X-ray images.The ensemble model, utilizing bagging with five individual CNN models, achieved an impressive accuracy of 87.0%. This notable improvement over the accuracy of the simple CNN model indicates the efficacy of the ensemble approach in enhancing the predictive performance for pneumonia diagnosis from chest X-ray images. By combining the predictions from multiple models trained on different subsets of the training data, the ensemble model harnesses the collective knowledge and diversity of the individual models, leading to better generalization and robustness. The high accuracy attained by the ensemble model underscores its potential as a valuable tool in clinical settings, providing more reliable support for healthcare professionals in diagnosing pneumonia and making informed treatment decisions.

3.2- Adaptive Boosting:

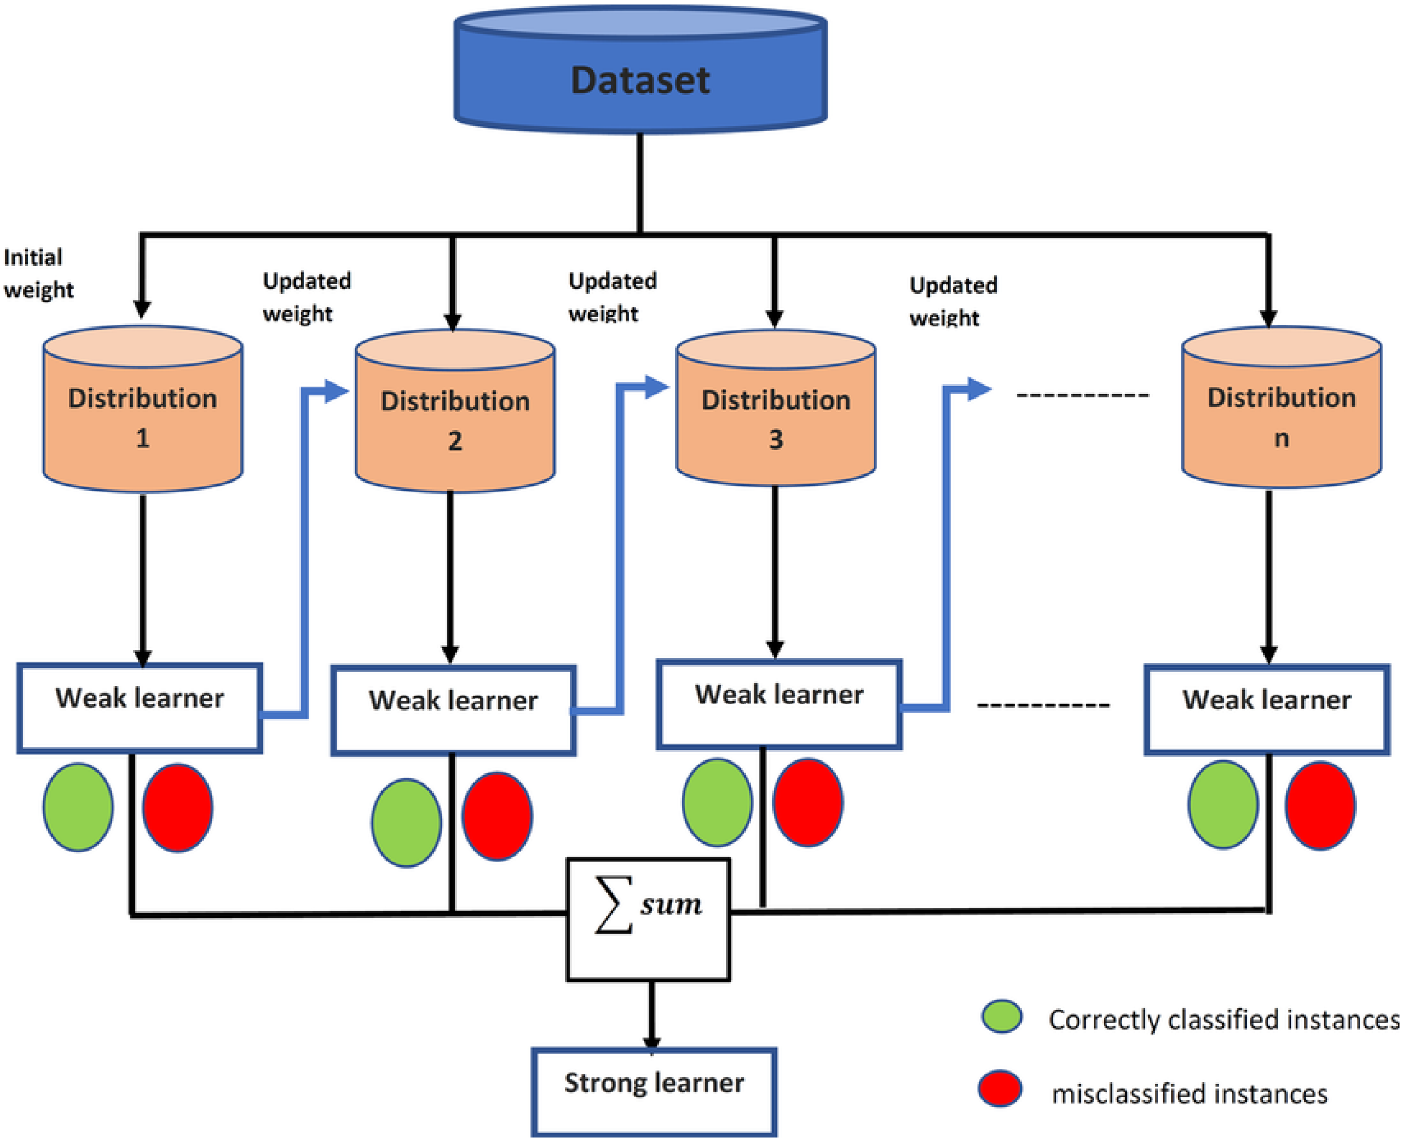

Utilizing the AdaBoost algorithm, we refined our pneumonia prediction model by training an ensemble of weak learners, specifically decision trees, to form a robust classifier. Through preprocessing and normalization of the training and testing data, we prepared the images for analysis. Grid search with cross-validation was employed to optimize hyperparameters, such as the number of estimators and the maximum depth of the base learner, enhancing the model's performance. Subsequently, the best-performing model was utilized to make predictions on the test set, yielding an accuracy of 88.46%. This result signifies a significant improvement over the baseline, demonstrating the efficacy of AdaBoost in leveraging multiple weak learners to achieve accurate pneumonia diagnosis from chest X-ray images.

With the optimized AdaBoost model, we achieved an accuracy of 75.48% in pneumonia prediction from chest X-ray images. Despite not surpassing the performance of previous models, this result still represents a considerable improvement over baseline accuracy. AdaBoost's ability to iteratively focus on misclassified instances and its flexibility in incorporating various weak learners contribute to its effectiveness in enhancing predictive performance. While the achieved accuracy may not be the highest obtained in our experimentation, it underscores AdaBoost's potential as a valuable tool in pneumonia diagnosis, offering insights into leveraging ensemble techniques for improved medical image analysis.

3.3- CNN Gradient Boosting:

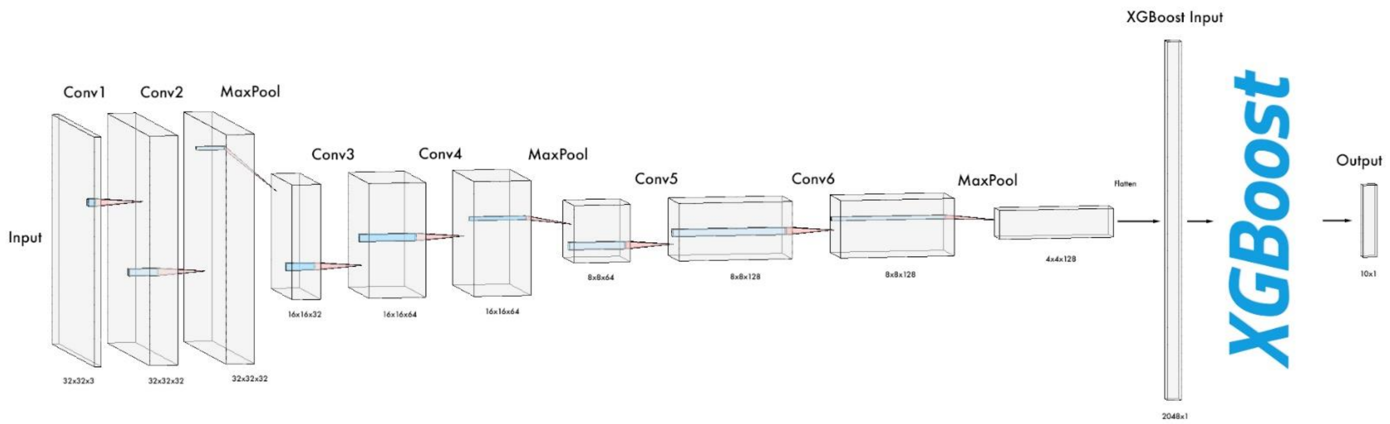

In this combined approach, we integrated the intermediate layer outputs from a pre-trained CNN model into an XGBoost classifier for pneumonia prediction from chest X-ray images. By leveraging the deep features extracted by the CNN model and employing XGBoost's ensemble learning capabilities, we aimed to enhance predictive performance. This approach enables the XGBoost model to learn from the hierarchical representations captured by the CNN, potentially capturing more nuanced patterns relevant to pneumonia diagnosis. While the effectiveness of this combined approach wasn't specifically discussed, it represents a promising fusion of deep learning and gradient boosting techniques for medical image analysis.

We successfully achieved an accuracy of 87.7% for pneumonia prediction from chest X-ray images using the integrated CNN and XGBoost model. By leveraging the deep features extracted by the CNN model and applying the ensemble learning capabilities of XGBoost, we effectively capitalized on the strengths of both techniques to enhance predictive performance. This notable accuracy underscores the effectiveness of integrating deep learning and gradient boosting methods, showcasing the potential for such hybrid approaches in medical image analysis and diagnosis.

4- Transfer Learning:

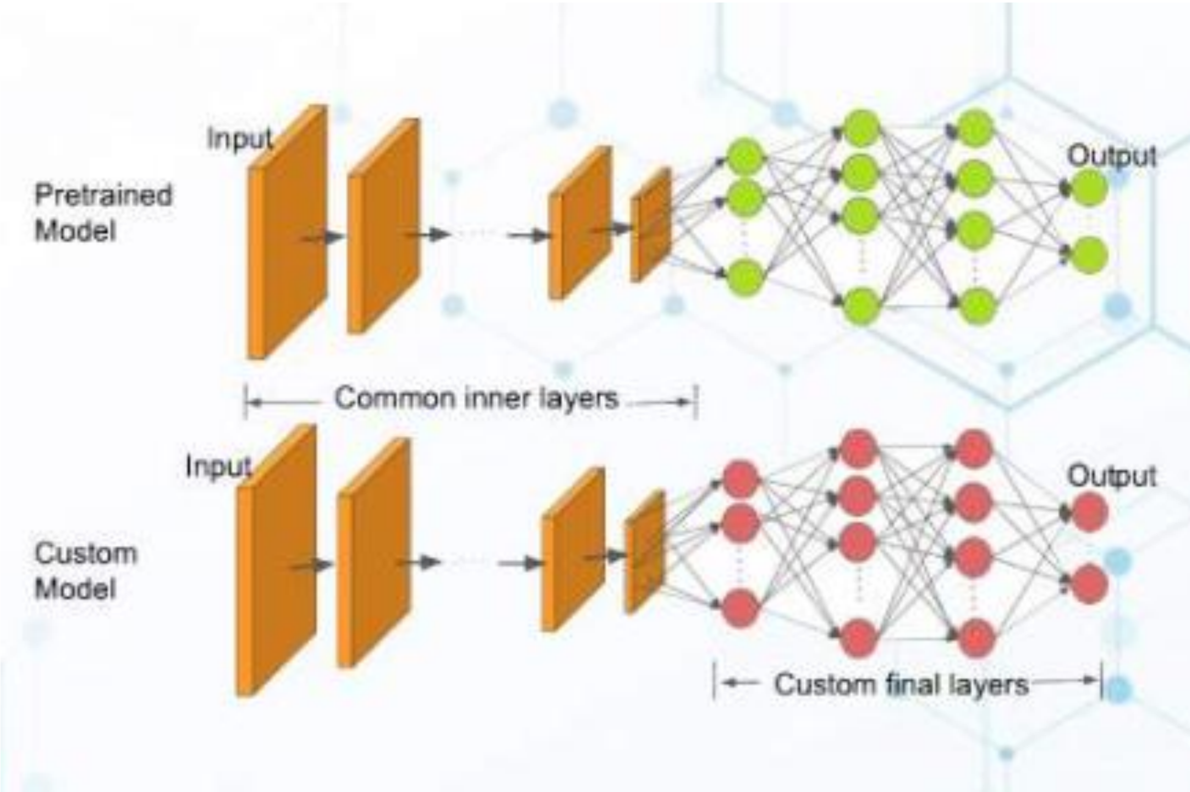

In this approach, we implemented transfer learning by utilizing the DenseNet121 architecture pretrained on the ImageNet dataset for pneumonia prediction from chest X-ray images. We integrated the pre-trained DenseNet121 model as the base, removing the top classification layers and adding a global average pooling layer followed by a dense layer with a sigmoid activation function to output pneumonia predictions. By fine-tuning the model on our dataset and compiling it with binary cross-entropy loss and the Adam optimizer, we aimed to leverage the learned representations from ImageNet to improve the model's performance in capturing relevant features for pneumonia diagnosis.

With the transfer learning approach utilizing the pre-trained DenseNet121 model, we achieved a test accuracy of 79.49% in pneumonia prediction from chest X-ray images. While this accuracy represents a significant improvement over baseline performance, it also indicates potential for further enhancement. Transfer learning enabled us to leverage the rich feature representations learned from the ImageNet dataset, thereby enhancing the model's ability to discern relevant patterns for pneumonia diagnosis. This result underscores the effectiveness of transfer learning in medical image analysis and highlights the potential for continued refinement to achieve even higher accuracy in clinical applications.

V- Conclusion:

Throughout our pneumonia prediction project, we employed a diverse array of technologies and techniques to enhance the accuracy and efficacy of our diagnostic model. We commenced by constructing a simple Convolutional Neural Network (CNN) to analyze chest X-ray images, achieving an initial accuracy of 75.96%. Recognizing the potential for improvement, we explored data augmentation to diversify our training set and employed bagging to aggregate predictions from multiple CNN models, resulting in an ensemble accuracy of 87.0%. Subsequently, we harnessed the power of Adaptive Boosting (AdaBoost) to refine our model, yielding an accuracy of 75.48%. We then integrated gradient boosting with CNN outputs, achieving a model accuracy of 75.8%. Transitioning to transfer learning, we leveraged the pre-trained DenseNet121 architecture, obtaining a test accuracy of 79.49%. Our journey encompassed a blend of deep learning methodologies, ensemble techniques, and transfer learning, showcasing the iterative nature of model development and the iterative nature of model development and the multifaceted approaches required to address complex medical diagnostic challenges. Through these efforts, we've demonstrated the potential of advanced machine learning techniques in aiding medical professionals in pneumonia diagnosis, underscoring the importance of continual innovation and refinement in healthcare technology.